

# Deep Graph Convolutional Networks

## Computer Vision Final Project

Nikola Janjušević

NYU

December 17th 2020

# Today's Papers

- [1] Francesca Pistilli et al. "Learning Graph-Convolutional Representations for Point Cloud Denoising". In: *The European Conference on Computer Vision (ECCV)*. 2020.
- [2] Martin Simonovsky and Nikos Komodakis. "Dynamic Edge-Conditioned Filters in Convolutional Neural Networks on Graphs". In: *CVPR*. 2017. URL:  
<https://arxiv.org/abs/1704.02901>.
- [3] Diego Valsesia, Giulia Fracastoro, and Enrico Magli. "Image Denoising with Graph-Convolutional Neural Networks". In: *2019 26th IEEE International Conference on Image Processing (ICIP)*. 2019.

# Table of Contents

1 Motivation

2 Convolution over Graphs

3 Graph Convolutional Denoising Network (GCDN)

# Motivation: graph representation of data I

- Irregular data domains: point clouds, sensor networks

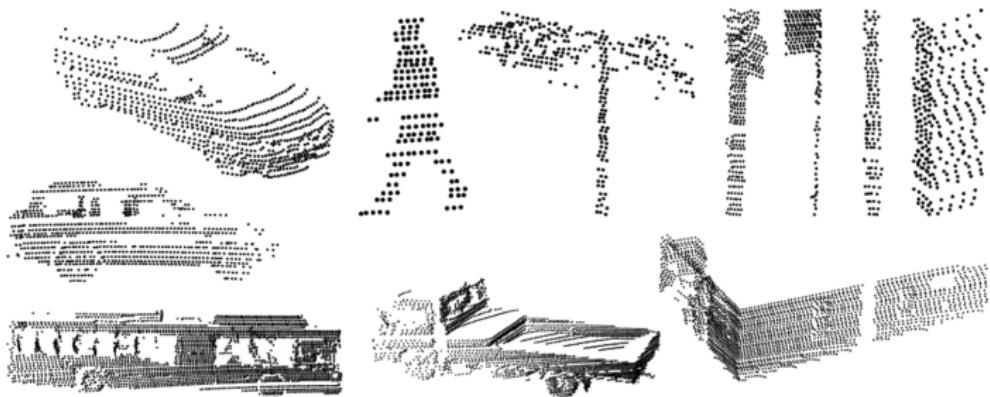


Figure: [2]

# Motivation: graph representation of data II

- Capture non-local relations in regular-domain data

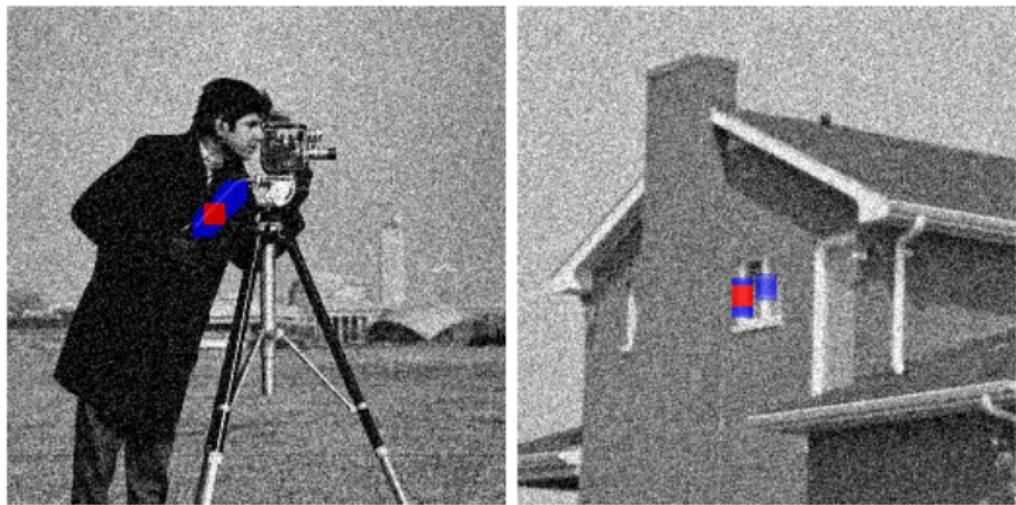


Figure: [BM3D]

# Table of Contents

1 Motivation

2 Convolution over Graphs

3 Graph Convolutional Denoising Network (GCDN)

## Linear Convolution as a special case:

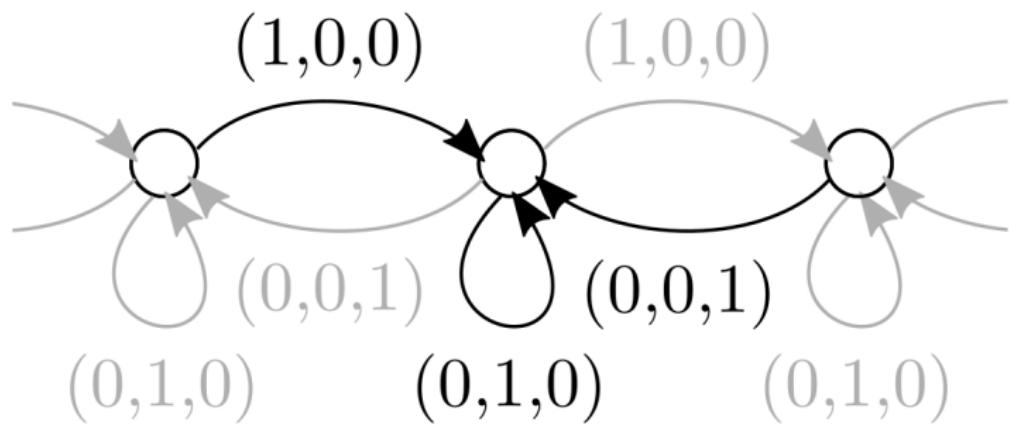


Figure: [2]

# Frequency and spatial domain formulations

- Notion of spectrum can be defined via Graph-Laplacian. Filtering then naturally arises.
- Building/applying Graph-Laplacian can be expensive, not nice for dynamically changing graphs.
- Spatial domain: define filtering as node/vertex signal aggregation. Less expensive but no clear formulation.

# Notation

Consider  $x \in \mathbb{R}^{CN}$ , a multi-channel signal with  $N$  pixels.

- $x[i] \in \mathbb{R}^C$ , feature vector of pixel  $i$ .
- $x^k \in \mathbb{R}^N$ , channel  $k$  of  $x$ .
- $\mathcal{G}(\mathcal{V}, \mathcal{E})$  directed graph with vertices  $\mathcal{V}$ , edges  $\mathcal{E}$ .
- $\mathcal{V}$  set of verticies,  $|\mathcal{V}| = N$
- $\mathcal{E}$  set of edges  $\mathcal{E} \subseteq \mathcal{V} \times \mathcal{V}$ .
- $\mathcal{N}(i) = \{j | (i, j) \in \mathcal{E}\}$  set of neighbors of vertex  $i$ .
- $\mathcal{L} : \mathcal{E} \rightarrow \mathbb{R}^D$  label map.
- $\mathcal{L}[i, j] \in \mathbb{R}^D$  label for edge connecting vertex  $j$  to  $i$ .

# Edge Conditioned Convolution I

## Definition (ECC [2])

Consider multi-channel signal  $x \in \mathbb{R}^{C_{in}N}$ , and *Label-Operator Mapping*  $\mathcal{F} : \mathbb{R}^M \rightarrow \mathbb{R}^{C_{out} \times C_{in}}$ . Define the *Edge-Conditioned Convolution* of  $x$  by  $\mathcal{F}$  as

$$y[i] = \sum_{j \in \mathcal{N}(i)} \mathcal{F}(\mathcal{L}[i, j])x[j] \quad (1)$$

$$= \sum_{j \in \mathcal{N}(i)} \mathbf{H}_{ij}x[j] \in \mathbb{R}^{C_{out}} \quad (2)$$

where we say  $\mathbf{H}_{ij}$  is the *edge-conditioned feature operator* from  $j$  to  $i$ .

- Consistent with multi-channel grid-convolution if  $\mathbf{H}_{ij}$  is stationary (only function of  $\tau = j - i$ ), though non-intuitive.

# Edge Conditioned Convolution II

- For 1D grid-convolution,  $\mathcal{F}(\mathcal{L}[i, j]) = \langle h, e_{j-i} \rangle$  ( $e_k$  dim  $k$  unit-vector), i.e. selecting coefficient  $j - i$  from filter  $h$ .
- When ordering of neighbors no-longer makes sense, stationary kernel can be abandoned.
- [2] learn label-operator mapping  $\mathcal{F}$  as 2-layer MLP
- Other spatial graph-convolution formulations don't have grid-convolution as special case, don't handle multi-channel signals properly, and/or use uninformative edge-labels
- [2] set SOA results on point-cloud classification

# Table of Contents

1 Motivation

2 Convolution over Graphs

3 Graph Convolutional Denoising Network (GCDN)

# GCDN

- Dynamically build K-Nearest Neighbors graph in feature-space to exploit non-local similarities for denoising images (like BM3D)
- Average local convolutions with non-local aggregations over graph (ECC)

# Motivating Results [3]

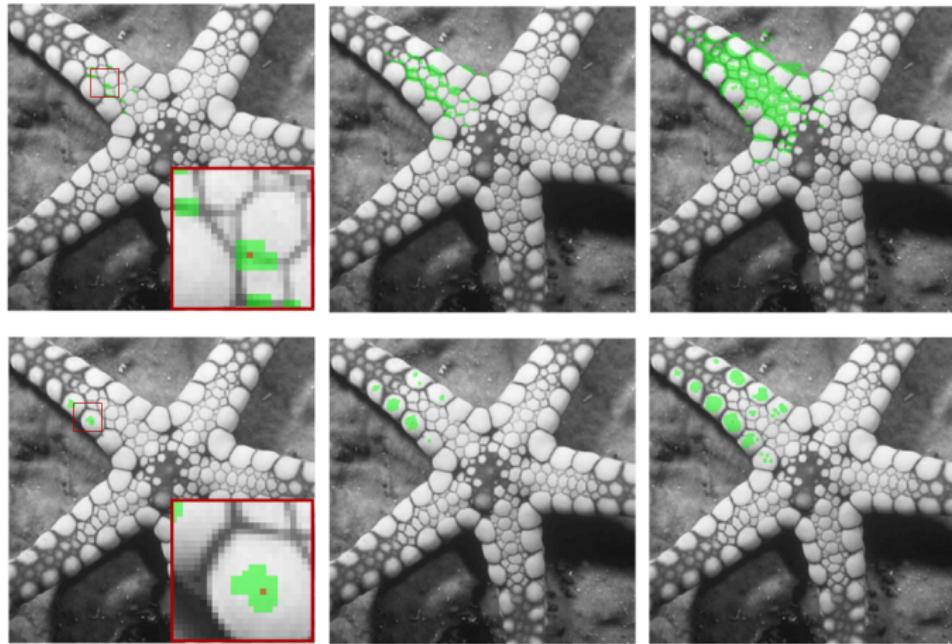
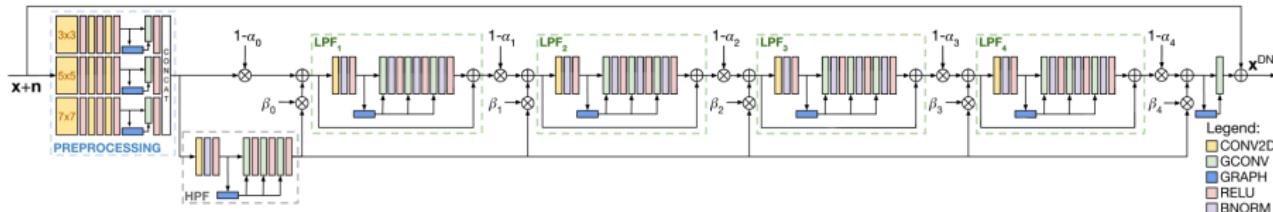


Fig. 6. Receptive field (green) of a single pixel (red) at the output of the three graph-convolutional layers in the  $\text{LPF}_1$  block with respect to the input of the first graph-convolutional layer in the block. Top row: gray pixel on an edge. Bottom row: white pixel in a uniform area.

# Architecture [3]

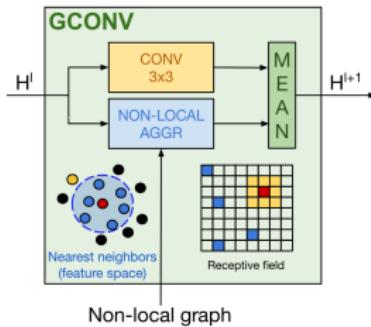


- Residual learning with batch-norm
- Build KNN graph for each “block” (not every convolution)
- All intermediate signals stay at same spatial resolution (no subsampling)
- Can be (loosely) viewed as unrolled proximal gradient descent with graph-laplacian operator,

$$\hat{x} = \arg \min \frac{1}{2} \| \mathbf{A}x - y \|_2^2 + \frac{\beta}{2} x^\top \mathbf{L}x$$

(for denoising  $\mathbf{A} = \mathbf{I}$ )

# GConv Module [3] |



- $$x^{(\ell+1)} = \frac{1}{2} \left( x^{(\ell+1),L} + x^{(\ell+1),NL} \right) + b^{\ell+1}$$
- $x^{(\ell+1),L}$  local 3x3 convolution

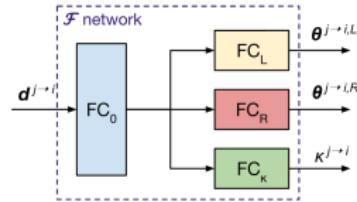
# GConv Module [3] II

- $x^{(\ell+1),NL}$  non-local ECC,

$$x^{(\ell+1),NL}[i] = \sum_{j \in \mathcal{N}^{(\ell)}(i)} \mathcal{F}^{(\ell)}(\mathcal{L}^{(\ell)}[i, j]) x^{(\ell)}[j] \quad (3)$$

$$= \sum_{j \in \mathcal{N}^{(\ell)}(i)} \mathbf{H}_{ij}^{(\ell)} x^{(\ell)}[j] \quad (4)$$

- Implement label operator map  $\mathcal{F}$  as low-rank MLP, via sum of  $r$  outer-products



# SOA Synthetic Denoising Results [3]

NATURAL IMAGE DENOISING RESULTS. METRICS ARE PNSR (dB) AND SSIM

Dataset	Noise $\sigma$	BM3D [5]	WNNM [31]	TNRD [34]	DnCNN [7]	$N^3$ Net [13]	NLRN [14]	GCDN
Set12	15	32.37 / 0.8952	32.70 / 0.8982	32.50 / 0.8958	32.86 / 0.9031	- / -	<b>33.16</b> / 0.9070 $\pm$ 1.06/0.0222	33.14 / <b>0.9072</b> $\pm$ 1.05/0.0214
	25	29.97 / 0.8504	30.28 / 0.8557	30.06 / 0.8512	30.44 / 0.8622	30.55 / -	<b>30.80</b> / <b>0.8689</b> $\pm$ 1.19/0.0305	30.78 / 0.8687 $\pm$ 1.20/0.0295
	50	26.72 / 0.7676	27.05 / 0.7775	26.81 / 0.7680	27.18 / 0.7829	27.43 / -	<b>27.64</b> / <b>0.7980</b> $\pm$ 1.31/0.0430	27.60 / 0.7957 $\pm$ 1.33/0.0443
BSD68	15	31.07 / 0.8717	31.37 / 0.8766	31.42 / 0.8769	31.73 / 0.8907	- / -	<b>31.88</b> / 0.8932 $\pm$ 2.48/0.0380	31.83 / <b>0.8933</b> $\pm$ 2.47/0.0383
	25	28.57 / 0.8013	28.83 / 0.8087	28.92 / 0.8093	29.23 / 0.8278	29.30 / -	<b>29.41</b> / 0.8331 $\pm$ 2.63/0.0564	29.35 / <b>0.8332</b> $\pm$ 2.66/0.0570
	50	25.62 / 0.6864	25.87 / 0.6982	25.97 / 0.6994	26.23 / 0.7189	26.39 / -	<b>26.47</b> / 0.7298 $\pm$ 2.78/0.0882	26.38 / <b>0.7389</b> $\pm$ 2.70/0.0877
Urban100	15	32.35 / 0.9220	32.97 / 0.9271	31.86 / 0.9031	32.68 / 0.9255	- / -	33.42 / 0.9348 $\pm$ 2.56/0.0324	<b>33.47</b> / <b>0.9358</b> $\pm$ 2.58/0.0326
	25	29.70 / 0.8777	30.39 / 0.8885	29.25 / 0.8473	29.97 / 0.8797	30.19 / -	30.88 / 0.9003 $\pm$ 2.63/0.0460	<b>30.95</b> / <b>0.9020</b> $\pm$ 2.65/0.0476
	50	25.95 / 0.7791	26.83 / 0.8047	25.88 / 0.7563	26.28 / 0.7874	26.82 / -	27.40 / <b>0.8244</b> $\pm$ 2.52/0.0711	<b>27.41</b> / 0.8160 $\pm$ 2.53/0.0758

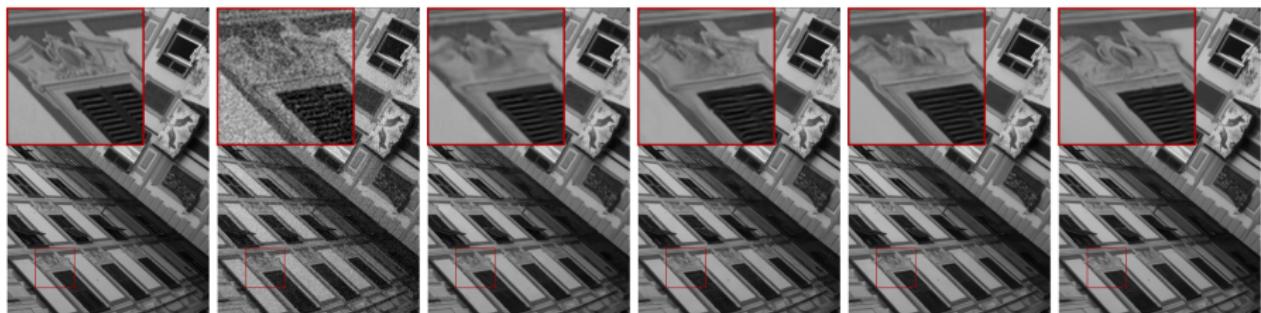


Fig. 11. Extract from Urban100 scene 13,  $\sigma = 25$ . Left to right: ground truth, noisy (20.16 dB), BM3D (30.40 dB), DnCNN (30.71 dB), NLRN (31.41 dB), GCDN (**31.53 dB**).

# SOA Real Image Denoising Results [3]

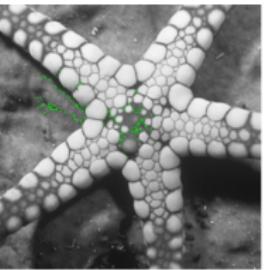
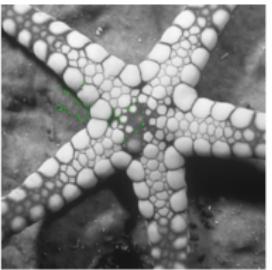
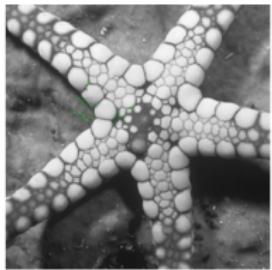
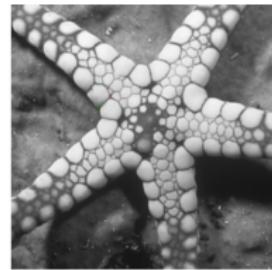
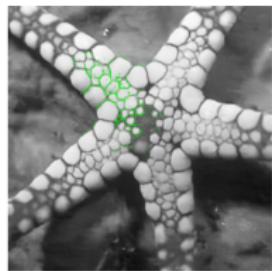
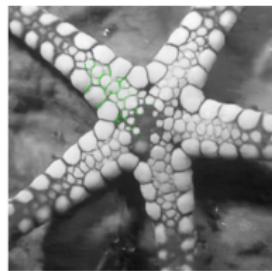
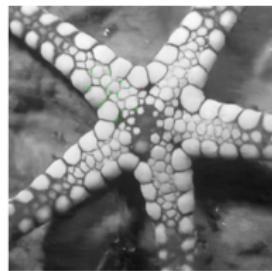
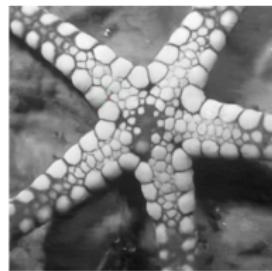
## REAL IMAGE DENOISING (SIDD DATASET)

	CBM3D	DnCNN	NLRN	GCDN
PSNR	$38.73 \pm 2.95$ dB	$39.98 \pm 3.17$ dB	$41.24 \pm 2.64$ dB	<b><math>41.48 \pm 2.15</math> dB</b>
SSIM	$0.9587 \pm 0.0138$	$0.9605 \pm 0.0158$	$0.9652 \pm 0.0144$	<b><math>0.9697 \pm 0.0132</math></b>



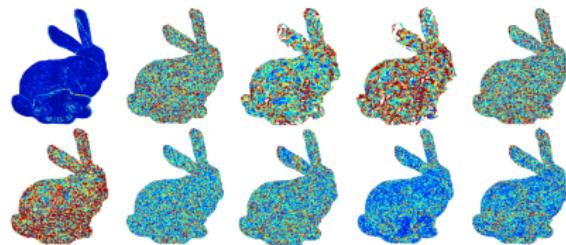
Fig. 13. Real image denoising. Left to right: ground truth, noisy (23.69 dB), CBM3D (36.90 dB), DnCNN (38.76 dB), NLRN (40.78 dB), GCDN (**41.21 dB**). Image intensity rescaled for visualization.

# Receptive Field: Learned feature-domain vs. image-domain

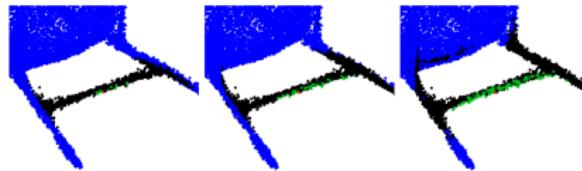


# Extensions

- SOA Point-Cloud denoising [1]



**Fig. 3.** Denoising results for  $\sigma = 0.015$ . Color represents unoriented normal angle error (red is high, blue is low). Top left to bottom right: clean point cloud (UNAE =  $3.75^\circ$ ), DGCNN ( $29.73^\circ$ ), APSS ( $26.29^\circ$ ), RIMLS ( $33.63^\circ$ ), AWLOP ( $29.18^\circ$ ), RPCA ( $37.50^\circ$ ), GLR ( $22.08^\circ$ ), PointCleanNet ( $23.63^\circ$ ), GPDNet MSE ( $16.62^\circ$ ), GPDNet MSE-SP ( $23.11^\circ$ ).



# References I

- [1] Francesca Pistilli et al. "Learning Graph-Convolutional Representations for Point Cloud Denoising". In: *The European Conference on Computer Vision (ECCV)*. 2020.
- [2] Martin Simonovsky and Nikos Komodakis. "Dynamic Edge-Conditioned Filters in Convolutional Neural Networks on Graphs". In: *CVPR*. 2017. URL:  
<https://arxiv.org/abs/1704.02901>.
- [3] Diego Valsesia, Giulia Fracastoro, and Enrico Magli. "Image Denoising with Graph-Convolutional Neural Networks". In: *2019 26th IEEE International Conference on Image Processing (ICIP)*. 2019.